

TestRIG: Randomized Testing of RISC-V CPUs using Direct Instruction Injection

Attending authors: Alexandre Joannou¹, Peter Rugg¹, Jonathan Woodruff¹, Franz A. Fuchs¹, Simon W. Moore¹

1 - University of Cambridge

peter.rugg@cl.cam.ac.uk

Project URL: www.github.com/CTSRD-CHERI/TestRIG



TestRIG

- TestRIG (Testing with Random Instruction Generation) is a testing framework for RISC-V implementations.
- We propose model-based verification: comparing execution between a design-under-test and a model implementation (e.g. Sail RISC-V).
- We use the RISC-V Formal Interface (RVFI) to observe the state changes after each instruction, and novel Direct Instruction Injection (DII) for test injection.
- Growing ecosystem: participants benefit from shared model and test-generation work by adopting the interface, regardless of implementation language.
- Integrated Sail RISC-V golden model, Piccolo, Flute, Toooba, RVBS (Bluespec), Ibex (Verilog), QEMU, and Spike (C).
- Standard interface libraries in a range of languages.

RISC-V Formal Interface (RVFI)

- RVFI is an existing trace format for formal verification using symbolic instructions.
- Exposes architectural signals such as the instruction encoding and memory and writeback information.
- An RVFI interface exports internal signals of an RTL design, or internal variables of a simulator or emulator.
- For more complex RTL designs, such as out-of-order microarchitectures, extracting the values may require preserving state for an RVFI report in a commit stage that could not previously access them.
- Extending the superscalar Toooba core for RVFI-DII required two new records in the Reorder Buffer.

```
struct RVFI_DII_Execution {  
    // ...  
    Bit#(32) rvfi_insn; // Executed instr  
    // Register writeback  
    Bit#(8) rvfi_rd_addr; // Written index  
    Bit#(64) rvfi_rd_wdata; // Written value  
    // Memory  
    Bit#(64) rvfi_mem_addr; // Address  
    Bit#(64) rvfi_mem_wdata; // Written data  
    Bit#(8) rvfi_mem_wmask; // Byte enables  
    Bit#(64) rvfi_mem_rdata; // Read data  
    // ...  
}
```

Listing 1. Definition of the RVFI type reported by RISC-V implementations. Not all fields are shown.

Direct Instruction Injection (DII)

- DII specifies the instructions expected in the output trace, separating instructions from memory addresses.
- Requires custom pipeline instrumentation, but enables greatly simplified sequence generation and shrinking.
- Must support resetting implementations between tests, including registers and memory.
- Implementations must rerun an instruction in the stream if it is flushed from the pipeline on mispredict.
- Interfaces provide a standard 8 MiB memory window.

```
struct RVFI_DII_Instruction {  
    // Instruction or reset command?  
    Bool rvfi_cmd;  
    // Time to inject instruction  
    Bit#(10) rvfi_time;  
    // Instruction word (32/16 bit)  
    Bit#(32) rvfi_insn;  
}
```

Listing 2. Definition of the DII type exposed to the implementations expressing the injected instruction.

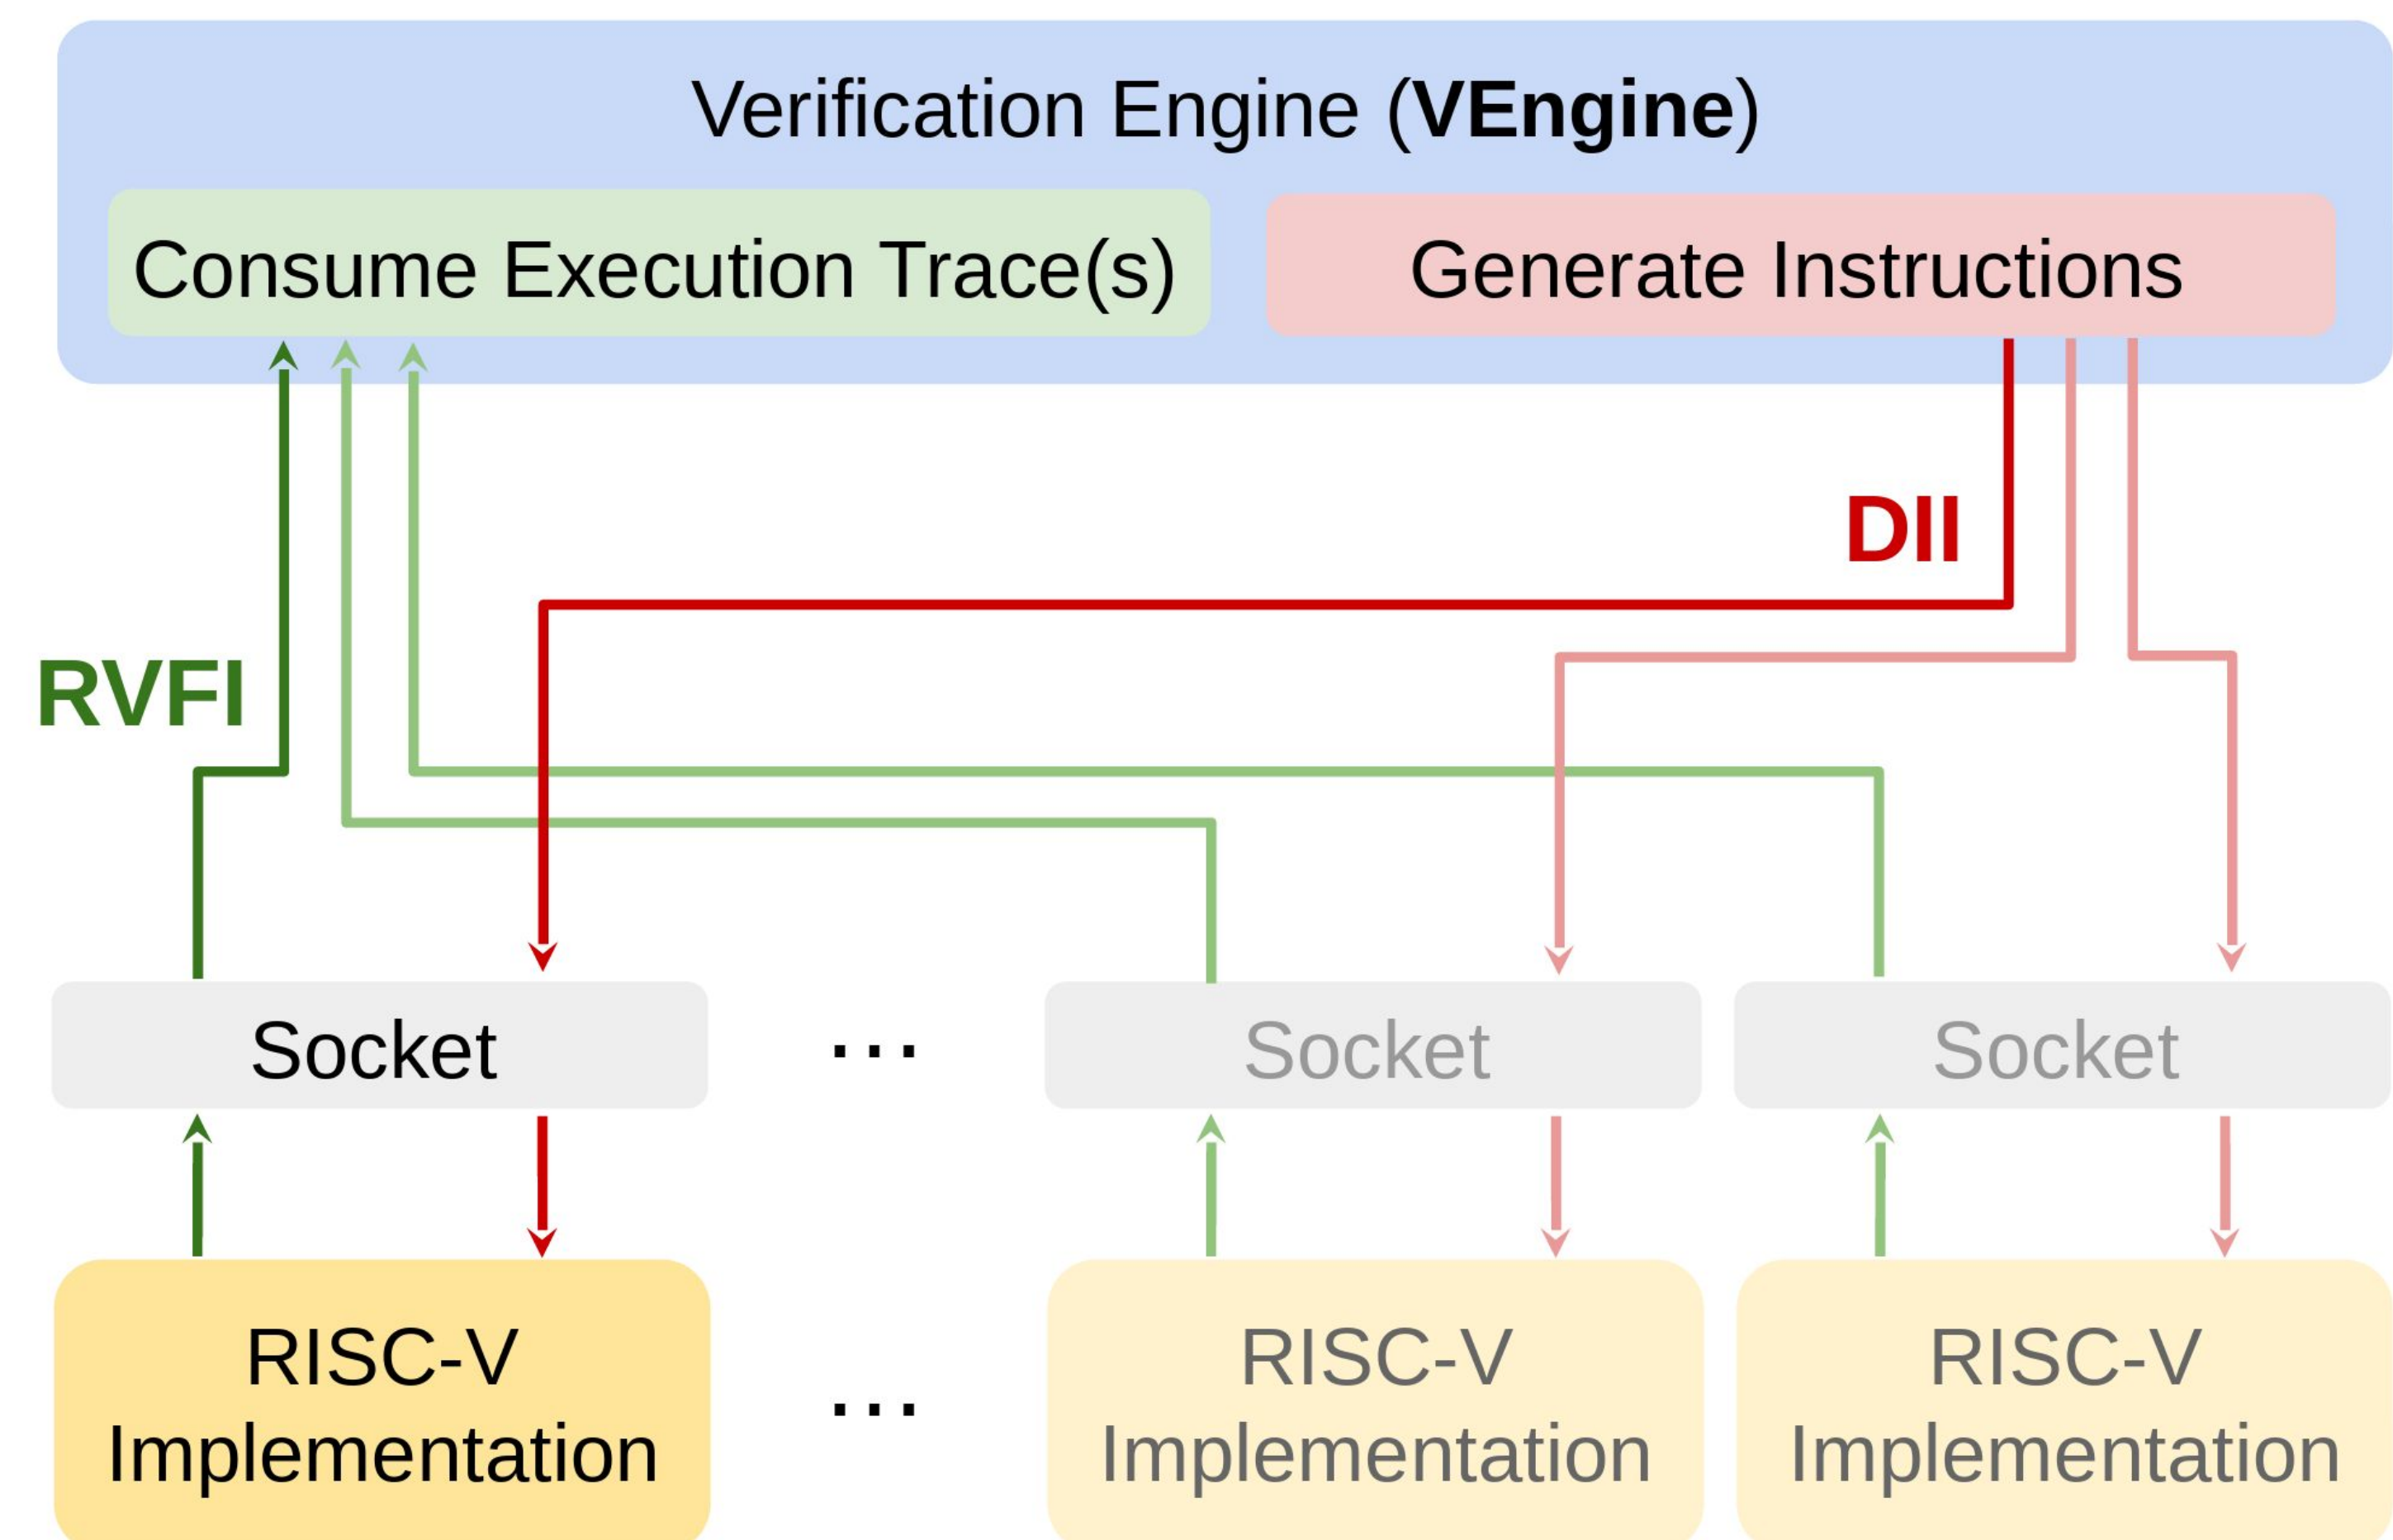


Figure 1. An illustration of the TestRIG ecosystem with a Verification Engine communicating with any two RISC-V implementations over sockets. The VEngine injects instruction sequences and compares the execution traces until it finds a divergence.

QuickCheck VEngine

- Leverages Haskell's *QuickCheck* library.
- Propositions passed into *QuickCheck* consist of a monadic function that issues instructions to both implementations and compares the reported results.
- *QuickCheck* attempts to disprove the proposition by generating test cases using our generator templates.
- On finding a counterexample, the trace is shrunk to find simpler test cases that still cause divergence.
- Trace serialisation enables a regression testing library without false positives when the specification is changed.
- Assertions ensure the intended behaviour is provoked.
- We have used QCVEngine to test our processors, including the CHERI RISC-V security extension.
- Eases per-instruction bringup.
- Found a security bug in capability compression code.
- It also accelerated debugging of a cache bug when switching Flute configurations (Listing 3).
- Enabled testing for speculative side-channels.

Intelligent shrinking

- DII allows traces to be automatically shrunk without control flow completely changing the test semantics.
- *QuickCheck*'s list shrinking function automatically deletes instructions unnecessary for divergence.
- Additional shrinking strategies include detecting *bypass* instructions that act as register moves.
- Library of candidate shrinks for each RISC-V instruction.
- Resulting counterexamples are much easier to understand, improving debugging experience.

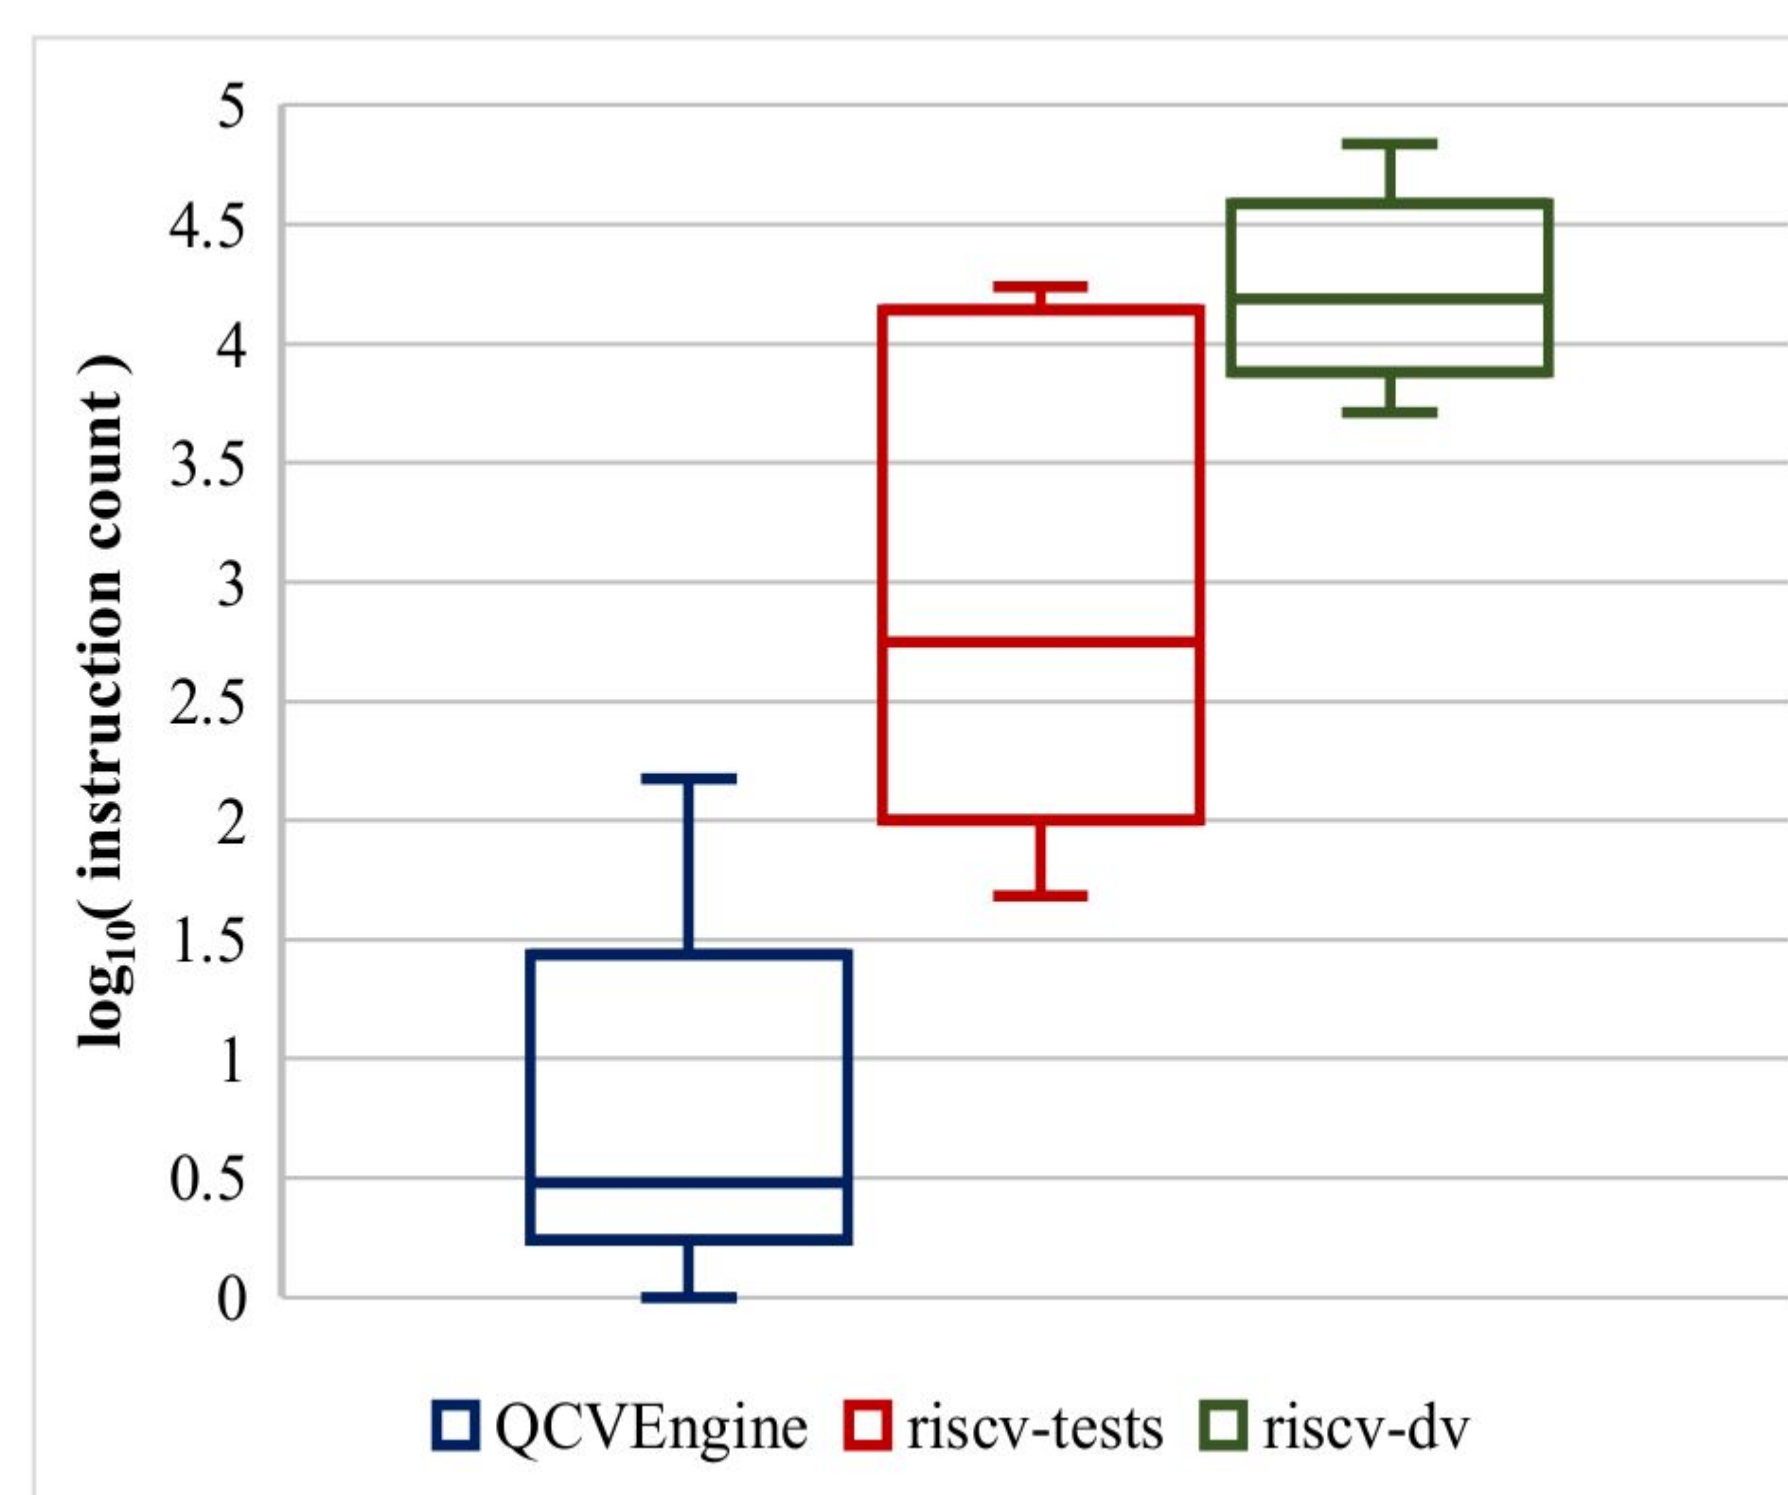


Figure 2. Size complexity of QCVEngine counterexamples for CHERI bringup compared to length of traditional tests.

```
lui x1, 0x40004  
slli x1, x1, 1  
lui x20, 0x40004 # Value used as data  
ori x3, x1, 1 # A page address  
lui x2, 0x40004  
slli x2, x2, 1 # The same page address  
lhu x4, x3[1] # Load from address  
sh x20, x2[2] # Store to same address  
lhu x2, x3[1] # Divergence on reloading
```

Listing 3. The shrunk counterexample produced by TestRIG to demonstrate a cache bug in the Flute processor: cache hits could return the line in the wrong way of a set.

Recursive templates

- Mechanism for directed-random test generation.
- Can specify static distribution of instructions.
- Use Haskell for *recursive templates* to reach *deep states* in the processor, e.g. setting up page-table mappings, switching privilege modes, etc.
- Compared to static tests, templates are general, allowing for varying addresses, register indices, and even sequences of instructions.
- They can also be composed arbitrarily.
- Our template library achieves good RISC-V architectural coverage (Sail RISC-V branch coverage), similar to the RISC-V test suite and RISC-V-DV.

```
surroundWithMemAccess x = random do  
    // Randomise fields  
    value <- bits 12 // 12 random bits  
    shift <- bits 6  
    // ...  
    return $ storeToAddress regAddr regData  
              offset value  
              shift  
              <> x  
              <> loadFromAddress regAddr  
                  offset regData
```

Listing 4. An example Haskell template to surround a given template with accesses to the same memory address.